

# CSC589 Introduction to Computer Vision

## Lecture 4

More on Histogram Equalization, Border  
Effect, Image Derivatives

Bei Xiao

# Last lecture

- Unsharp masking
- Gaussian Filter
- Image histograms
- Basic image tutorial with Python

# Separability example

2D convolution  
(center location only)

$$\begin{bmatrix} 1 & 2 & 1 \\ 2 & 4 & 2 \\ 1 & 2 & 1 \end{bmatrix} * \begin{bmatrix} 2 & 3 & 3 \\ 3 & 5 & 5 \\ 4 & 4 & 6 \end{bmatrix}$$

The filter factors  
into a product of 1D  
filters:

$$\begin{bmatrix} 1 & 2 & 1 \\ 2 & 4 & 2 \\ 1 & 2 & 1 \end{bmatrix} = \begin{bmatrix} 1 \\ 2 \\ 1 \end{bmatrix} \times \begin{bmatrix} 1 & 2 & 1 \end{bmatrix}$$

Perform convolution  
along rows:

$$\begin{bmatrix} 1 & 2 & 1 \end{bmatrix} * \begin{bmatrix} 2 & 3 & 3 \\ 3 & 5 & 5 \\ 4 & 4 & 6 \end{bmatrix} = \begin{bmatrix} & 11 & \\ & 18 & \\ & 18 & \end{bmatrix}$$

Followed by convolution  
along the remaining column:


# Matrix Multiplication

$$\mathbf{u} \otimes \mathbf{v} = \mathbf{u}\mathbf{v}^T = \begin{bmatrix} u_1 \\ u_2 \\ u_3 \\ u_4 \end{bmatrix} \begin{bmatrix} v_1 & v_2 & v_3 \end{bmatrix} = \begin{bmatrix} u_1 v_1 & u_1 v_2 & u_1 v_3 \\ u_2 v_1 & u_2 v_2 & u_2 v_3 \\ u_3 v_1 & u_3 v_2 & u_3 v_3 \\ u_4 v_1 & u_4 v_2 & u_4 v_3 \end{bmatrix}.$$


# Today's Lecture

- More on image histograms
- Border effect and padding
- Image Derivatives/gradients

# Image Histogram Equalization



Notice the “shape” of the histogram is preserved!



Same number of pixels  
In each bin!

# Questions

1. Why don't we do the equalization directly on the histograms? What will happen to the image if it has a perfectly flat histograms?
2. What is CDF? Why are there steps (zig-zags) on the CDF?
3. How to make sure the “shape” of the PDF is preserved in the transformation?

# Let's look at in more details

We start with a gray-scale jpeg image of 8 by 8

52	55	61	66	70	61	64	73
63	59	55	90	109	85	69	72
62	59	68	113	144	104	66	73
63	58	71	122	154	106	70	69
67	61	68	104	126	88	68	70
79	65	60	70	77	68	58	75
85	71	64	59	55	61	65	83
87	79	69	68	65	76	78	94

[https://en.wikipedia.org/wiki/Histogram\\_equalization](https://en.wikipedia.org/wiki/Histogram_equalization)



# Let's look at in more details

The histograms for this image (8 by 8 pixels) is shown in the following table (intensity that has zero pixels are skipped):

Value	Count	Value	Count	Value	Count	Value	Count	Value	Count
52	1	64	2	72	1	85	2	113	1
55	3	65	3	73	2	87	1	122	1
58	2	66	2	75	1	88	1	126	1
59	3	67	1	76	1	90	1	144	1
60	1	68	5	77	1	94	1	154	1
61	4	69	3	78	1	104	2		
62	1	70	4	79	2	106	1		
63	2	71	2	83	1	109	1		

# Let's look at in more details

The cumulative distribution function (CDF) is shown below

Value	cdf	scaled cdf
52	1	0
55	4	
58	6	
59	9	
60	10	
61	14	
62	15	
...	...	
154	64	255

This cdf shows that the min value in the subimage is 52 and max is 154. The cdf of value 154 corresponding to the total number of pixels (64)

Number of pixels

# Let's look at in more details

How do we compute the normalized CDF?

$$h(v) = \text{round} \left( \frac{cdf(v) - cdf_{min}}{(M \times N) - cdf_{min}} \times (L - 1) \right)$$

Cdf(v): original cdf of pixel v

Cdfmin: minimum non-zero value of the cdf

M×N: number of pixels, e.g. 64 (8x8)

L: 256

# Quiz:

How do we compute normalized CDF of pixel 62?

Value	cdf	scaled cdf
52	1	0
55	4	
58	6	
59	9	
60	10	
61	14	
62	15	?
...	...	
154	64	255

$$h(v) = \text{round} \left( \frac{cdf(v) - cdf_{min}}{(M \times N) - cdf_{min}} \times (L - 1) \right)$$

# Quiz:

How do we compute normalized CDF of pixel 62?

Value	cdf	scaled cdf
52	1	0
55	4	
58	6	
59	9	
60	10	
61	14	
62	15	?
...	...	
154	64	255

$$\begin{aligned} H(62) &= \text{round}((15-1)/ \\ &63*255) \\ &= 57 \end{aligned}$$

# Let's look at in more details

The cumulative distribution function (CDF) is shown below

Value	cdf	scaled cdf
52	1	0
55	4	12
58	6	22
59	9	32
60	10	36
61	14	53
62	15	57
...	...	...
154	64	255


# Let's look at in more details

Now we can directly map the scaled cdf back to pixel values using the look up table we derived above

	52	55	61	66					
	↗	↗	↗	↗					
[	0	12	53	93	146	53	73	166	]
	65	32	12	215	235	202	130	158	
	57	32	117	239	251	227	93	166	
	65	20	154	243	255	231	146	130	
	97	53	117	227	247	210	117	146	
	190	85	36	146	178	117	20	170	
	202	154	73	32	12	53	85	194	
	206	190	130	117	85	174	182	219	

# Let's look at in more details

Notice that the minimum value (52) is now 0 and the maximum value (154) is now 255.




Original

Equalized



# Let's look at in more details

Notice that the minimum value (52) is now 0 and the maximum value (154) is now 255.



Original

Equalized

# Homework 2: histogram equalization

Hint:

#Compute cdf in Python:

```
np.histogram(im.flatten(),nbr_bins,normed=True)
```

```
cdf = imhist.cumsum()
```

```
# normalize
```


```
cdf = 255 * cdf / cdf[-1]
```

```
# Using linear interpretation of cdf to find new pixels.
```

```
im2 = np.interp(im.flatten(),bins[:-1],cdf)
```

# Back to Convolution

- Let's blur a flat gray image:



How should it look like?

# Back to Convolution

- Convolve with a box kernel

1/9	1/9	1/9				
1/9	1/9	1/9				
	20	20		20	20	
1/9	1/9	1/9				
	20	20		20	20	
	20	20	20	20		
	20	20	20	20		

# Border Handling

Depending on how you do the convolution, you could end up with 3 different images.



266x266 Image

full



256x256 Image

same



246x246 Image

valid

# Back to Convolution

- Zero padding:

Filled in borders with zeros, computed everywhere the kernel touches.

`Numpy.convolve(a,v,mode = 'full' or 'same')`

This returns the convolution at each point of overlap, with an output shape of  $(N+M-1,)$ . At the end-points of the convolution, the signals do not overlap completely, and boundary effects may be seen.

# Back to Convolution

- Convolve with a box kernel, suppose there are zeros outside the image matrix.

$\frac{1}{9}$	$\frac{1}{9}$	$\frac{1}{9}$					
$\frac{1}{9}$	$\frac{1}{9}$	$\frac{1}{9}$					
	20	20		20		20	
$\frac{1}{9}$	$\frac{1}{9}$	$\frac{1}{9}$					
	20	20		20		20	
	20	20		20		20	
	20	20		20		20	
	20	20		20		20	




8.9	13.3	13.3	8.9
13.3	20.0	20.0	13.3
13.3	20.0	20.0	13.3
8.9	13.3	13.3	8.9

# Back to Convolution

- This is called “same” in MATLAB/NumPy



4 × 4



4 × 4



# Back to Convolution

1/9	1/9	1/9			
1/9	1/9	1/9			
1/9	1/9	1/9			
		20	20	20	20
		20	20	20	20
		20	20	20	20
		20	20	20	20




2.2	4.4	6.6	6.6	4.4	2.2
4.4	8.9	13.3	13.3	8.9	4.4
6.6	13.3	20.0	20.0	13.3	6.7
6.6	13.3	20.0	20.0	13.3	6.7
4.4	8.9	13.3	13.3	8.9	4.4
2.2	4.4	6.6	6.6	4.4	2.2

# Back to Convolution

- This is called “full” in MATLAB/NumPy.



4 × 4



6 × 6

# Back to Convolution

- Only compute at places where the kernel fits the image

$\frac{1}{9}$ 20	$\frac{1}{9}$ 20	$\frac{1}{9}$ 20		20
$\frac{1}{9}$ 20	$\frac{1}{9}$ 20	$\frac{1}{9}$ 20		20
$\frac{1}{9}$ 20	$\frac{1}{9}$ 20	$\frac{1}{9}$ 20		20
20	20	20	20	



20	20
20	20

# Back to Convolution

- This is called “valid” in MATLAB/NumPy.




$4 \times 4$



$2 \times 2$

# To summarize

- Python:
  - `Numpy.convolve(a,v,mode)`
  - `Scipy.signal.convolve2d(a,v,mode, boundary, value)`



# There are other methods

- The first two methods that I described fill missing values in by substituting zero
- Can fill in values with different methods
  - Reflect image along border
  - Pull values from other side
- Read Chapter 3.2 “padding”.

# There are other methods

zero





valid



Nearest neighbor



# Edge detection




- Convert a 2D image into a set of curves
  - Extracts salient features of the scene
  - More compact than pixels

Source: N. Snavely




# Origin of Edges



- Edges are caused by a variety of factors
- It is still a very active research topic! We will learn more about it!

# Image derivate and edges

- An edge is a place of rapid change in image intensity function



# Image derivatives

- How can we differentiate a *digital* image  $F[x,y]$ ?
  - Option 1: reconstruct a continuous image,  $f$ , then compute the derivative
  - Option 2: take discrete derivative (finite difference)  $\frac{\partial f}{\partial x}[x,y] \approx F[x+1,y] - F[x,y]$

How would you implement this as a linear filter?

$$\frac{\partial f}{\partial x} : \begin{array}{|c|c|c|} \hline & & \\ \hline & & \\ \hline & & \\ \hline \end{array}$$

$H_x$




$$\frac{\partial f}{\partial y} : \begin{array}{|c|c|c|} \hline & & \\ \hline & & \\ \hline & & \\ \hline \end{array}$$

$H_y$

# Image gradient

- The *gradient* of an image:  $\nabla f = \left[ \frac{\partial f}{\partial x}, \frac{\partial f}{\partial y} \right]$

The gradient points in the direction of most rapid increase in intensity



The **edge strength** is given by the gradient magnitude:

$$\|\nabla f\| = \sqrt{\left(\frac{\partial f}{\partial x}\right)^2 + \left(\frac{\partial f}{\partial y}\right)^2}$$

The gradient direction is given by:


$$\theta = \tan^{-1} \left( \frac{\partial f}{\partial y} / \frac{\partial f}{\partial x} \right)$$

- how does this relate to the direction of the edge?

# Image gradient




$f$



$\frac{\partial f}{\partial x}$





$\|\nabla f\| = \sqrt{\left(\frac{\partial f}{\partial x}\right)^2 + \left(\frac{\partial f}{\partial y}\right)^2}$



$\frac{\partial f}{\partial y}$

# Gaussian Filter

- Gaussian = normal distribution function




$$K(i, j) = \frac{1}{Z} \exp \left( -\frac{i^2 + j^2}{2\sigma^2} \right)$$

# Derivative of Gaussian

- Take the derivative of the filter with respect to  $i$ :



$$\frac{\partial K(i, j)}{\partial i} = \frac{-i}{\sigma^2 Z} \exp\left(-\frac{i^2 + j^2}{2\sigma^2}\right)$$

- Filter looks like:




- Basically blur then take the derivative

# Derivative of Gaussian






# 2D edge detection filters



Gaussian


$$h_{\sigma}(u, v) = \frac{1}{2\pi\sigma^2} e^{-\frac{u^2+v^2}{2\sigma^2}}$$




derivative of Gaussian (x)

$$\frac{\partial}{\partial x} h_{\sigma}(u, v)$$



# Derivative of Gaussian filter



x-direction



y-direction



# The Sobel operator

- Common approximation of derivative of Gaussian

$$\frac{1}{8} \begin{array}{|c|c|c|} \hline -1 & 0 & 1 \\ \hline -2 & 0 & 2 \\ \hline -1 & 0 & 1 \\ \hline \end{array}$$


$s_x$

$$\frac{1}{8} \begin{array}{|c|c|c|} \hline 1 & 2 & 1 \\ \hline 0 & 0 & 0 \\ \hline -1 & -2 & -1 \\ \hline \end{array}$$

$s_y$

- The standard defn. of the Sobel operator omits the  $1/8$  term
  - doesn't make a difference for edge detection
  - the  $1/8$  term **is** needed to get the right gradient value

# Example of Sobel filtered image



# Take-home exercise: Create image derivatives using Sobel filter

1. Generate an image of a rotated rectangle

```
import numpy as np
from scipy import ndimage
import matplotlib.pyplot as plt
```

```
im = np.zeros((256, 256))
im[64:-64, 64:-64] = 1
im = ndimage.rotate(im, 15, mode='constant')
```

2. Blur the image using a Gaussian filter

```
im = ndimage.gaussian_filter(im, 8)
```

3. Apply Sobel filter to both x and y direction.

```
sx = ndimage.sobel(im, axis=0, mode='constant')
```

4. Display the original image, x-derivatives, y-derivatives, and the gradient magnitude. You can use `np.hypot` to compute magnitude.

See here: <http://docs.scipy.org/doc/numpy/reference/generated/numpy.hypot.html>)

# Next class

- Laplacian of Gaussian
- Steerable filter
- Fourier transform
- Reading: Chapter 3.1-3.3! Very important to keep up the reading.
- You should have read Chapter 1. Skip Chapter 2 since we haven't covered it.